

Post-Cleaning Verification

The dataset now has a new column: `loyalty_tier` — this wasn't present in the original analysis. The source file has been updated. All 1,050 rows have a valid loyalty tier with no missing values, suggesting it was added as part of a deliberate data enrichment.

Record Flow: Raw → Cleaned

 Cleaning funnel ▼	
Collapsed. Click to open chart.	
Stage	Rows
Raw data	1,050
After deduplication (–50 duplicates)	1,000
After invalid row removal (–67 rows)	933

The 67 removed rows break down as: 18 invalid emails, 20 invalid ages, and 21 amounts outside the 0–\$10,000 range (some rows overlap multiple checks). The final cleaned dataset contains 933 valid, unique orders.

New Column — Loyalty Tier

 Loyalty tier distribution 

Collapsed. Click to open chart.

Loyalty Tier	Orders
Bronze	272 (26%)
Platinum	267 (25%)
Silver	264 (25%)
Gold	247 (24%)

The loyalty tier split is remarkably even — all four tiers (Bronze, Silver, Gold, Platinum) hold roughly 25% of orders each. This near-perfect balance suggests the tiers were algorithmically assigned or randomly distributed, and may warrant investigation to confirm whether they reflect genuine customer behaviour or a synthetic assignment.

After normalising casing and removing the raw fragmentation, the six countries consolidate cleanly: Brazil (156), Canada (148), Nigeria (146), Australia (144), Vietnam (137), India (113). The distribution is fairly even — no single market dominates, which may indicate this is a globally distributed customer base.

Monthly Trend (Cleaned)

 Cleaned monthly trend 

Collapsed. Click to open chart.

With outlier amounts removed, monthly revenue now falls consistently in the \$1,600–\$4,700 range, which is a far more credible picture than the raw data's wild spikes (\$10M

- in April and September 2025). Order volumes are steady at 7–16 orders/month with no strong seasonality signal over the 20-month window.

The key action item from this check: the `loyalty_tier` column is new and should be added to the pipeline's cleaning scope — specifically, the  `deduped` node's column list and any downstream schema references should be updated to include it. Would you like me to update the pipeline to handle it?